


Slides Production: Prof. Yau Kok Lim

Dr. Saad Aslam

Office: Room AE-327, 3rd Floor, NUB

Contact: +603 74918622
 (Ext: 7143)

Email: saada@sunway.edu.my




Computer Networking: A Top-down Approach, 7th edition.
 Jim Kurose, Keith Ross
 Pearson

Section	Topic	Slides
4.1	Introduction	2
	• Per-router Control Plane	3
	• Logically Centralized Control Plane	4
4.3	The Internet Protocol (IP): Forwarding and Addressing in the Internet	
	• IPv4 Addressing	5-7
	• How does a host get IP address?	8-10
	• Network Address Translation (NAT)	11-12
	• IPv6: Transitioning from IPv4 to IPv6	13-15
5.2	Routing Algorithms: Overview	16-17
	• The Link-State (LS) Routing Algorithm	18-21
	• The Distance-Vector (DV) Routing Algorithm	22-27
5.3	Hierarchical Routing	28
5.3	Routing Information Protocol (RIP)	29
5.4	Border Gateway Protocol (BGP)	30-31



4.1 Introduction


Sender Host



- **Network Layer**


- Provide logical communication between two hosts
 - Sender host: Encapsulate segments into datagrams
 - Intermediate host (Router):
 - Examine datagram header
 - Perform routing and forwarding
 - Destination host: Pass datagrams to transport layer
- Two planes
 - **Data plane**
 - Move packet from router's input to a router's output
 - Local function
 - **Control plane**
 - Determine a route from source to destination (routing)
 - Network-wide function
 - Two approaches
 - **Per-router control plane**
 - Each router perform routing separately to compute forwarding table
 - **Logically centralized control plane**
 - A centralized remote controller interact with local control agent in routers to compute forwarding table

4.1 Per-router Control Plane



- A traditional approach implemented in all routers.
- Each router computes its own forwarding table in a distributed manner.

4.1 Logically Centralized Control Plane (also called Software-defined Networking, SDN)



- A newer approach implemented in a remote controller.
- Each router contains a forwarding table that is computed and distributed by a remote controller.




4.3 IPv4 Addressing

- IPv4 address
 - 32 bit
 - 2^{32} possible IP addresses
 - Problem
 - 2^{32} IP addresses is used up
 - Solution
 - Use IPv6
 - 128 bit
- Each IPv4 address identify an interface (connection between a host/router and a link)
 - Each IP address is globally unique
 - Router has multiple interface
 - E.g.: This router has 3 interfaces
 - Each interface has an IP address


223.1.1.1 = 11011111 00000001 00000001 00000001

223 1 1 1



4.3 IPv4 Addressing


- IPv4 address
 - High-order bits
 - Indicate subnet
 - Low-order bits
 - Indicate host
- Each subnet
 - Hosts have similar high-order bits
 - Hosts are physically relevant to each other (e.g.: an organization)
 - E.g.: This router has 3 subnet
 - 223.1.1.0/24
 - /24 is subnet mask
 - Indicate 24 high-order bits
 - Has 3 host interface
 - 223.1.1.1
 - 223.1.1.2
 - 223.1.1.3
 - Has 1 router interface
 - 223.1.1.4
 - 223.1.3.0/24
 - 223.1.2.0/24



4.3 IPv4 Addressing


- How many subnets? 6 subnet

- 223.1.1.0/24
- 223.1.2.0/24
- 223.1.3.0/24
- 223.1.7.0/24
- 223.1.8.0/24
- 223.1.9.0/24




- Hard-coded by system admin in a file
 - Windows
 - control-panel->network->configuration->tcp/ip->properties
 - UNIX
 - /etc/rc.config
- Dynamic Host Configuration Protocol (DHCP)
 - Allows a host to dynamically obtain an IP address from a server when it joins a network
 - “plug-and-play”

4.3 How does a host get IP address? Dynamic Host Configuration Protocol




4.3 How does a host get IP address? Dynamic Host Configuration Protocol



- Network address translation (NAT)
 - local network uses just one IP address as far as outside world is concerned
 - Advantages
 - range of addresses not needed from ISP
 - just one IP address for all devices
 - can change addresses of devices in local network without notifying outside world
 - can change ISP without changing addresses of devices in local network
 - devices inside local network not explicitly addressable, visible by outside world (a security plus)


4.3 Network Address Translation (NAT)




4.3 IPv6: Transitioning from IPv4 to IPv6

- IPv6


- not all routers can be upgraded simultaneously
- How to provide downward compatible (compatible with IPv4)?
 - **Tunneling**
 - Carry IPv6 datagram as payload in IPv4 datagram among IPv4 routers




Logical view:



Physical view:




4.3 IPv6: Transitioning from IPv4 to IPv6



- IPv6 adoption
 - Google: 8% of clients access services via IPv6
 - Takes long time (up to 20 years) for adoption

5.2 Routing Protocols: Overview

- Routing algorithm
 - Centralized or Decentralized
 - **Centralized**
 - Also called Link-State Algorithm
 - Router has complete network-wide information including
 - Network topology (Routers and Links)
 - Link cost
 - **Decentralized / Distributed**
 - Also called Distance Vector Algorithm
 - Router has neighbor nodes' information including
 - Network topology
 - Link cost
 - Neighbor nodes exchange information (i.e. route cost) among themselves
 - Static or Dynamic
 - **Static**
 - Route change slowly
 - Route change due to changes in network topology and link cost
 - **Dynamic**
 - Route change faster
 - Router update link changes frequently



- Network representation
 - Network, $G = (N, E)$
 - Routers, $N = \{u, v, w, x, y, z\}$
 - Links,
 $E = \{(u, v), (u, w), (u, x), (v, w), (v, x), (w, x), (w, y), (w, z), (x, y), (y, z)\}$
 - Link cost, $c(x, x')$
 - E.g.: $c(u, w) = 5$
 - Related to network condition, such as bandwidth, network congestion level, delay
 - Cost of path $(x_1, x_2, x_3, \dots, x_p)$


$$= c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$$
 - E.g.: Cost of path (u, v, w, z)

$$= c(u, v) + c(v, w) + c(w, z)$$

$$= 2 + 3 + 5 = 10$$
- Routing algorithm
 - Calculate least-cost path (route) between u and z




5.2 The Link-State (LS) Routing Algorithm



- Routing representation
 - $c(x, x')$: Link cost
 - $D(v)$: Current cost of path from source to destination v
 - $p(v)$: Previous node (neighbor of node v) along the current least-cost path from source to node v
 - N' : Subset of nodes. Node v is in N' if the least-cost path from the source to v is known
- Dijkstra's algorithm
 - Determine the least cost path (route) between u and z
 - Node (Source)
 - Calculate least-cost path (route) from itself to all other nodes in the network
 - Use least-cost path to update its forwarding table
 - After k iterations, node (source) know least-cost paths to k destination node


5.2 The Link-State (LS) Routing Algorithm: An Example



- Routing representation
 - $c(x, x')$: Link cost
 - $D(v)$: Current cost of path from source to destination v
 - $p(v)$: Previous node (neighbor of node v) along the current least-cost path from source to node v
 - N' : Subset of nodes. Node v is in N' if the least-cost path from the source to v is known
- Dijkstra's algorithm

Step	N'	$D(v), p(v)$	$D(w), p(w)$	$D(x), p(x)$	$D(y), p(y)$	$D(z), p(z)$
0	u	2,u	5,u	1,u	∞	∞
1	ux	2,u	4,x		2,x	∞
2	uxy	2,u	3,y		4,y	
3	uxyv		3,y		4,y	
4	uxyvw					4,y
5	uxyvwz					


5.2 The Link-State (LS) Routing Algorithm: An Example



- Routing representation
 - $c(x, x')$: Link cost
 - $D(v)$: Current cost of path from source to destination v
 - $p(v)$: Previous node (neighbor of node v) along the current least-cost path from source to node v
 - N' : Subset of nodes. Node v is in N' if the least-cost path from the source to v is known

Dijkstra's algorithm outcome:

- Least-cost path tree



Dijkstra's algorithm outcome:


- Forwarding table for node u

Destination	Link
v	(u, v)
x	(u, x)
y	(u, x)
w	(u, x)
z	(u, x)




5.2 The Link-State (LS) Routing Algorithm: Problem


- Problem
 - Oscillation
 - Happen when link cost is related to network congestion level (or traffic load)
 - Example: Node *B*, *C* and *D* send to node *A*




Path for
 $B: (B, A)$
 $C: (C, B, A)$
 $D: (D, A)$



given these costs,
 find new routing....
 resulting in new costs
 Path for
 $B: (B, C, D, A)$
 $C: (C, D, A)$
 $D: (D, A)$




given these costs,
 find new routing....
 resulting in new costs
 Path for
 $B: (B, A)$
 $C: (C, B, A)$
 $D: (D, C, B, A)$



given these costs,
 find new routing....
 resulting in new costs
 Path for
 $B: (B, C, D, A)$
 $C: (C, D, A)$
 $D: (D, A)$

- Solution
 - Randomize the time neighbor nodes exchange link cost among themselves

5.2 The Distance-Vector (DV) Routing Algorithm




- Routing representation
 - $c(x, x')$: Link cost
 - $D_x(y)$: Distance vector
 - Cost of the least-cost path from node x and node y
 - v : Node x 's neighbor nodes

- Bellman-Ford equation:

$$D_x(y) = \min_v \{c(x, v) + D_v(y)\} \text{ for each node } y \text{ in } N$$

- \min_v means that, among node x 's neighbor nodes, node x choose a neighbor node with minimum $c(x, v) + D_v(y)$
- Neighbor nodes exchange information (i.e. distance vector) among themselves
 - Node x
 - Calculate $D_x(y)$
 - Send $D_x(y)$ to all neighbor nodes v only if there is changes on $D_x(y)$
 - Each neighbor node receive $D_x(y)$
 - Keep track and update the distance vector $D_x(y)$ using Bellman-Ford equation

5.2 The Distance-Vector (DV) Routing Algorithm



- Routing representation
 - $c(x, x')$: Link cost
 - $D_x(y)$: Distance vector
 - Cost of the least-cost path from node x and node y
 - v : Node x 's neighbor nodes
 - Bellman-Ford equation

$$D_x(y) = \min_v \{c(x, v) + D_v(y)\}$$

- Example on Bellman-Ford equation:
 - Calculate least-cost path (route) between u and z
 - $D_v(z) = 5, D_w(z) = 3, D_x(z) = 3$
 - $c(u, v) = 2, c(u, w) = 5, c(u, x) = 1$
 - $D_u(z) = \min \{c(u, v)+D_v(z), c(u, w)+D_w(z), c(u, x)+D_x(z)\}$
 $= \min \{7, 8, 4\}$
 $= 4$

5.2 The Distance-Vector (DV) Routing Algorithm: An Example

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

**node x
table**

	x	y	z
x	0	2	7
y	∞	∞	∞
z	∞	∞	∞

	x	y	z
x	0	2	3
y	2	0	1
z	7	1	0


**node y
table**

	x	y	z
x	∞	∞	∞
y	2	0	1
z	∞	∞	∞

**node z
table**

	x	y	z
x	∞	∞	∞
y	∞	∞	∞
z	7	1	0

time



5.2 The Distance-Vector (DV) Routing Algorithm: An Example

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x table

	x	y	z
from	0	2	7
x	0	2	7
y	∞	∞	∞
z	∞	∞	∞

node y table

	x	y	z
from	∞	∞	∞
x	∞	∞	∞
y	2	0	1
z	∞	∞	∞

node z table

	x	y	z
from	∞	∞	∞
x	∞	∞	∞
y	∞	∞	∞
z	7	1	0


	x	y	z
from	0	2	3
x	0	2	3
y	2	0	1
z	7	1	0

	x	y	z
from	0	2	3
x	0	2	3
y	2	0	1
z	3	1	0

	x	y	z
from	0	2	3
x	0	2	3
y	2	0	1
z	3	1	0

	x	y	z
from	0	2	3
x	0	2	3
y	2	0	1
z	3	1	0

time




5.2 The Distance-Vector (DV) Routing Algorithm: Problem and Solution

- Problem

- **Good news** travel **Fast**, **Bad news** travel **Slow**


- *Good news travel fast*

- t_0 : y detect link-cost change from 4 to 1
- t_1 : z receive $D_y(x) = 1$ from y , so z update $D_z(x) = 2$
- t_2 : y receive $D_z(x) = 2$ from z
 - Since $D_y(x) < c(y, z) + D_z(x)$
 - No update for $D_y(x)$ at z




- *Bad news travel slow*

- t_0 : y detect link-cost change from 4 to 60
- Since $D_y(x) > c(y, z) + D_z(x)$
 - $D_y(x) = 60$; $c(y, z) + D_z(x) = 1 + 5 = 6$
 - y update its own $D_y(x) = 6$
- t_1 : z receive $D_y(x) = 6$ from y , so z update $D_z(x) = 7$
- t_2 : y receive $D_z(x) = 7$ from z , so y update $D_y(x) = 8$
- t_3 : z receive $D_y(x) = 8$ from y , so z update $D_z(x) = 9$
- **Routing loop** between y and z until $D_z(x) > 50$
 - So, routing loop has 44 iteration!




5.2 The Distance-Vector (DV) Routing Algorithm: Problem and Solution

- Solution
 - **Poisoned Reverse**
 - Solve “*Bad news travel slow*”
 - How it works?
 - Since Z route through Y to get to X
 - Z tell Y the distance to X is infinite $D_Z(x) = \infty$
 - So, Y do not route to X via Z
 - t_0 : y detect link-cost change from 4 to 60
 - Since $D_y(x) > c(y, z) + D_z(x)$
 - $D_y(x) = 60$; $c(y, z) + D_z(x) = 1 + 5 = 6$
 - y update its own $D_y(x) = 6$
 - t_1 : z receive $D_y(x) = 6$ from y
 - Using poisoned reverse, z tell lie that $D_z(x) = \infty$
 - z route to x directly and update its $D_z(x) = 50$
 - t_2 : y receive $D_z(x) = \infty$ from z , so y update $D_y(x) = 60$
 - t_3 : z receive $D_y(x) = 60$ from y , so z update $D_z(x) = 50$
 - t_4 : y receive $D_z(x) = 50$ from y , so y update $D_y(x) = 51$



5.3 Hierarchical Routing

- Group routers into **Autonomous Systems (AS)**
 - Advantage
 - Scalability
 - Reduce routing info as the number of routers increase
 - Administrative autonomy
 - A single company can own and manage AS
 - Figure show three AS
 - AS1, AS2, AS3
 - Figure show AS1 has four router
 - 1a, 1b, 1c, 1d
- Hierarchical Routing
 - Two types
 - **Intra-AS routing**
 - Forward packet within AS
 - Example
 - Routing Information Protocol (RIP)
 - **Inter-AS routing**
 - Forward packet between AS
 - Example
 - Border Gateway Protocol (BGP)




- **Gateway router**
 - Characteristic
 - Located at the edge of AS
 - Has link to router in the other AS
 - Used by inter-AS routing to forward packet between AS
 - Figure show three gateway router
 - 1b, 1c, 2a, 3a

5.3 Routing Information Protocol (RIP)



- RIP
 - Distance vector algorithm
 - $D_x(y)$: Distance vector
 - Cost of the least-cost path from node x and node y
 - $D_x(y) = \text{Number of hops}$
 - Maximum $D_x(y)$ is 15 hops
 - $c(x,x')$: Link cost
 - $c(x,x') = 1$
 - Node exchange **RIP Response Message** with neighbors every 30s
 - If no RIP Response Message heard after 180s, it indicates a link failure
 - Route Recovery
 - New RIP Response Message sent to neighbour if routing table changes, so link failure is propagated to entire network quickly
 - Poison reverse prevent loops
 - Note:
 - In the Figure
 - Node indicate router
 - Router A, B, C, D
 - Link indicate subnet
 - Subnet u, v, w, x, y, z




From router A to destination **subnets**:

subnet	hops
u	1
v	2
w	2
x	3
y	3
z	2

5.4 Border Gateway Protocol (BGP)

- **eBGP**
 - Send subnet reachability info to neighboring AS
 - Enable subnet to advertise its existence to the rest of the internet: “I am here”
- **iBGP**
 - Distribute subnet reachability information to all routers within an AS.
 - Determine least-cost path to other networks based on subnet reachability information and company policy.
 - Subnet reachability information: (see next slide)
 - Company policy:
 - E.g.: An AS charges high rate, so choose another path
- All routers throughout the network always learn new prefix and create new entry for the prefix in their respective forwarding table



5.4 Border Gateway Protocol (BGP)

- eBGP and iBGP service
 - **Step 1:** Router 3a use eBGP session between 3a and 1c to send subnet reachability info to AS1
 - **Step 2:** Router 1c use iBGP session to distribute prefix to all routers within AS1
 - **Step 3:** Router 1b use eBGP session between 1b and 2a to re-advertise subnet reachability info to AS2

- Subnet reachability info contain
 - **Prefix**
 - Example: AS3 can reach subnets
 - 128.16.64/24
 - 128.16.65/24
 - 128.16.66/24
 - Then, AS3 send prefix 128.16.64/22 to AS1
- **AS-PATH**
 - Contain AS through which prefix has passed
 - Example: Router 1b send AS-PATH = { AS3 AS1 } to router 2a
- **NEXT-HOP**
 - Contain IP address of a router's interface within an AS which prefix has passed
 - Example: Router 3a send NEXT-HOP = {IP address of router 3a's interface} to router 1c

